

# 1 Ideas

## 1.1 Line search

Line search in ?? is not working very well.

line search

$$r' = \underset{r \in [0, 1]}{\operatorname{argmin}} \quad \underset{\tilde{D}_{KL}(r)^\uparrow}{KL(q^t + r(s - q^t) \| p)} = \nabla_r f(\tilde{q}_t)$$

grad descent,

$$r \leftarrow r - \eta \nabla_r \tilde{D}_{KL}(r)$$

func. grad,

$$\nabla_{\tilde{q}_t} f(\tilde{q}_t)$$

$$\nabla_r KL(q^t + r(s - q^t) \| p)$$

$$= \log \frac{s}{p}$$

$$\int \underbrace{(q^t + r(s - q^t))}_{\tilde{q}_t} \log \left( \frac{q^t + r(s - q^t)}{p} \right) dr$$

$$-g_n = \langle \nabla f(n), s \rangle - \langle \nabla f(n), n \rangle$$

$$= \nabla_r \int \tilde{q}_t \log \frac{\tilde{q}_t}{p}$$

$$= \int \left( \nabla_r \tilde{q}_t \right) \log \frac{\tilde{q}_t}{p} + \cancel{\tilde{q}_t \nabla_r \log \frac{\tilde{q}_t}{p}}$$

$\downarrow$   
 $\frac{1}{\tilde{q}_t} \nabla_r \tilde{q}_t$   
 $\downarrow$   
 $\frac{1}{s - q^t}$

$$= \int (s - q^t) \log \frac{\tilde{q}_t}{p} + (s - q^t)$$

$$= \int (s - q^t) \left( 1 + \log \frac{\tilde{q}_t}{p} \right)$$

$$= \int s \log \frac{\tilde{q}_t}{p} - \int q^t \log \frac{\tilde{q}_t}{p}$$

$$= \langle s, \log \frac{\tilde{q}_t}{p} \rangle - \langle q^t, \log \frac{\tilde{q}_t}{p} \rangle$$

$$= \mathbb{E}_s [\log \tilde{q}_t - \log p] - \mathbb{E}_{q^t} [\log \tilde{q}_t - \log p]$$

$$= \mathbb{E}_s [\nabla_{\tilde{q}_t} f(\tilde{q}_t)] - \mathbb{E}_{q^t} [\nabla_{q^t} f(\tilde{q}_t)]$$

```

1 def line_search_dkl(weights, locs, diags, mu_s, cov_s, x, k):
2     """Perform line search for the best step size gamma.
3
4     Uses gradient ascent to find gamma that minimizes
5     KL(q_t + gamma (s - q_t) || p)
6
7     Args:
8         weights: weights of mixture components of q_t
9         locs: means of mixture components of q_t
10        diags: deviations of mixture components of q_t
11        mu_s: mean for LMO Solution s
12        cov_s: cov matrix for LMO solution s
13        x: target distribution p
14        k: iteration number of Frank-Wolfe
15    Returns:
16        Computed gamma
17    """
18    def softmax(v):
19        return np.log(1 + np.exp(v))
20    # no. of samples to approximate  $\nabla_\gamma$ 
21    N_samples = 10
22    # Create current iter  $q_t$ 
23    weights = [weights]
24    qt_comps = [
25        Normal(
26            loc=tf.convert_to_tensor(locs[i]),
27            scale=tf.convert_to_tensor(diags[i])) for i in range(len(locs))
28    ]
29    qt = Mixture(
30        cat=Categorical(probs=tf.convert_to_tensor(weights)),
31        components=qt_comps,
32        sample_shape=N)
33    qt = InfiniteMixtureScipy(stats.multivariate_normal)
34    qt.weights = weights[0]
35    qt.params = list(
36        zip([[l] for l in locs], [[softmax(np.dot(d, d))] for d in diags]))
37    # samples from  $q_t$ 
38    sample_q = qt.sample_n(N_samples)
39    # create and sample from s
40    s = stats.multivariate_normal([mu_s],
41                                  np.dot(np.array([cov_s]), np.array([cov_s])))
42    sample_s = s.rvs(N_samples)
43    #  $q_{t+1}$  is mixture of  $q_t$  and s with weights  $(1-\gamma)$  and  $\gamma$ 
44    # Set its corresponding parameters and weights
45    new_locs = copy.copy(locs)
46    new_diags = copy.copy(diags)
47    new_locs.append([mu_s])
48    new_diags.append([cov_s])
49    # initialize  $\gamma$ 
50    gamma = 2. / (k + 2.)
51    # no. steps of gradient ascent
52    n_steps = 10
53    prog_bar = ed.util.Progbar(n_steps)
54    for it in range(n_steps):
55        print("line_search iter %d, %.5f" % (it, gamma))
56        new_weights = copy.copy(weights)
57        new_weights[0] = [(1. - gamma) * w for w in new_weights[0]]
58        new_weights[0].append(gamma)
59        # create  $q_{t+1}^\gamma$ 
60        q_next = InfiniteMixtureScipy(stats.multivariate_normal)
61        q_next.weights = new_weights[0]
62        q_next.params = list(
63            zip([[l] for l in new_locs], [[np.dot(d, d)] for d in new_diags]))
64        # Computes  $\mathbb{E}[\dots] \propto \sum_v \log p - \log q_{t+1}^\gamma$ 
65        def px_qx_ratio_log_prob(v):
66            Lambda = 1.

```

```

67     ret = x.log_prob([v]).eval()[0] - q_next.log_prob(v)
68     ret /= Lambda
69     return ret
70     # Samples w.r.t s
71     rez_s = [
72         px_qx_ratio_log_prob(sample_s[ss]) for ss in range(len(sample_s))
73     ]
74     # Samples w.r.t qt+1
75     rez_q = [
76         px_qx_ratio_log_prob(sample_q[ss]) for ss in range(len(sample_q))
77     ]
78     # Gradient ascent step, step size decreasing as  $\frac{1}{it+1}$ 
79     gamma = gamma + 0.1 * (sum(rez_s) - sum(rez_q)) / (N_samples *
80         (it + 1.))
81     # Projecting it back to [0, 1], too small range?
82     # FIXME(sauravshekhar) if projected to 0, all iterations will be same?
83     if gamma >= 1 or gamma <= 0:
84         gamma = max(min(gamma, 1.), 0.)
85     break
86     return gamma

```

changes for measuring variance of  $\mathbb{E}_s[\cdot]$  and  $\mathbb{E}_{q_{t+1}^\gamma}[\cdot]$ .

```

1     ...
2     grad_gamma = []
3     for it in range(n_steps):
4         ...
5         # Samples w.r.t s
6         rez_s = np.asarray([
7             px_qx_ratio_log_prob(sample_s[ss]) for ss in range(len(sample_s))
8         ])
9         # Samples w.r.t qt+1
10        rez_q = np.asarray([
11            px_qx_ratio_log_prob(sample_q[ss]) for ss in range(len(sample_q))
12        ])
13        grad_gamma.append({'E_s': rez_s, 'E_q': rez_q, 'gamma': gamma})
14        ...
15        # Write grad_gamma to outdir/line_search_samples_<n_samples>.npy.<fw_iter>

```

Metrics on original version.




Figure 1: gamma with iterations for different n\_samples

## 1.2 Adaptive Frank-Wolfe from smoothness estimators

Algorithm 1 of [Pedregosa et al., 2018]

Code changes.




Figure 2:  $E_s$  with different  $n\_samples$




Figure 3:  $E_q$  with different  $n\_samples$

?? begins with iter 1 as iter 0 has very high variance

### 1.3 Measuring smoothness

in progress

Computing optimal  $\gamma$  directly from eqn 1 of [Pedregosa et al., 2018]

$$f(\mathbf{x}_{t+1}) \leq f(\mathbf{x}_t) + \gamma \langle \nabla f(\mathbf{x}_t), \mathbf{s}_t - \mathbf{x}_t \rangle + \frac{\gamma^2}{2} L_t \|\mathbf{s}_t - \mathbf{x}_t\|^2$$

$$\Rightarrow L_t \geq \frac{f(\mathbf{x}_{t+1}) - f(\mathbf{x}_t) + \gamma \langle \nabla f(\mathbf{x}_t), \mathbf{s}_t - \mathbf{x}_t \rangle}{\frac{\gamma^2}{2} \|\mathbf{s}_t - \mathbf{x}_t\|^2}$$

$$\frac{f(\mathbf{x}_{t+1}) - f(\mathbf{x}_t) + \gamma \langle \nabla f(\mathbf{x}_t), \mathbf{s}_t - \mathbf{x}_t \rangle}{\frac{\gamma^2}{2} \text{KL}(\mathbf{s}_t || \mathbf{q}_t)}$$

Code changes.

```
1 def kl(mu_q, std_q, mu_p, std_p):
2     ...
3
4 def func_dot_product():
5     ...
```

### 1.4 Other optimization algorithm

todo

As shown in link ??, Frank-Wolfe converges slower than Projected Gradient Descent in Practice. See [Locatello et al., 2017] to see why we use FW and if it can be replaced. (will have to derive new convergence proofs and boosting won't be as integrated into the optimization algorithm as before).

---

**Algorithm 1:** Adaptive Frank-Wolfe for Boosting BBVI

---

**Input:**  $q_0 \in \mathcal{D}$ , initial Lipschitz estimate  $L_{-1}$ , line search parameters  $\tau > 1, \eta \in (0, 1]$

```
1 for  $t = 1 \dots T$  do
2    $s_t \leftarrow LMO_{\mathcal{A}}(\nabla f(q^t))$ 
3    $g_t \leftarrow \langle \nabla f(q_t), q_t \rangle - \langle \nabla f(q_t), s_t \rangle$  //  $\text{Gap} \geq 0$ 
4   Find smallest integer  $i$  s.t
5    $f(q_t + \gamma_t(s_t - q_t)) \leq Q_t(\gamma_t)$  Where
6    $Q_t(\gamma) := f(q_t) - \gamma g_t + \frac{\gamma^2 L_t}{2} d(s_t, q_t)$  // Quadratic upper bound ??
7    $L_t \leftarrow \tau^i \eta L_{t-1}$  and  $\gamma_t \leftarrow \min\left(\frac{g_t}{L_t d(s_t, q_t)}, 1\right)$ 
8 end
9 return  $q_T$ 
```

---

## 1.5 Entropy Regularization and Noise addition using Optimal Transport

In LMO, [Locatello et al., 2018] uses Entropy Regularization in place of norm constrained optimization. It can be replaced with something simpler See [Tolstikhin et al., 2017] [Dong Liu, 2018] [Bernton et al., 2017] [Jordan et al., 1998] [here](#) And [Peyré et al., 2017] part 4.

## 1.6 Port code to Tensorflow Probability

see ?? . Issue is if `tfp.edward2` will have support for Variational Inference and ELBO etc.

## References

- [Bernton et al., 2017] Bernton, E., Jacob, P. E., Gerber, M., and Robert, C. P. (2017). Inference in generative models using the wasserstein distance. *arXiv preprint arXiv:1701.05146*.
- [Dong Liu, 2018] Dong Liu, Minh Thnh Vu, S. C. L. K. R. (2018). Entropy-regularized optimal transport generative models. *arXiv preprint arXiv:1811.06763*.
- [Jordan et al., 1998] Jordan, R., Kinderlehrer, D., and Otto, F. (1998). The variational formulation of the fokker–planck equation. *SIAM journal on mathematical analysis*, 29(1):1–17.
- [Locatello et al., 2018] Locatello, F., Dresdner, G., Khanna, R., Valera, I., and Rätsch, G. (2018). Boosting black box variational inference. *arXiv preprint arXiv:1806.02185*.
- [Locatello et al., 2017] Locatello, F., Khanna, R., Ghosh, J., and Rätsch, G. (2017). Boosting variational inference: an optimization perspective. *arXiv preprint arXiv:1708.01733*.
- [Pedregosa et al., 2018] Pedregosa, F., Askari, A., Negiar, G., and Jaggi, M. (2018). Step-size adaptivity in projection-free optimization. *arXiv preprint arXiv:1806.05123*.
- [Peyré et al., 2017] Peyré, G., Cuturi, M., et al. (2017). Computational optimal transport. Technical report, École Normale Supérieure.
- [Tolstikhin et al., 2017] Tolstikhin, I., Bousquet, O., Gelly, S., and Schoelkopf, B. (2017). Wasserstein auto-encoders. *arXiv preprint arXiv:1711.01558*.